

Two models, because a realtime agent cannot wait

The Talker speaks. The Thinker plans between turns and never touches the line.



Every claim resolves to a lookup



Deepgram nova-3 in, aura-2 out · GPT-5.6 Luna both sides · prompts are files, swappable without a deploy